Growing A Tetris Player with Big Data
Project Report by Group 22

Li Jiaxin Cassandra
A0131652N

Loh Han Tong, Victor
A0135808B

Ryan Tan Wen Jun
A0135747X

Tan Yu Wei
A0142255M

April 21, 2018

1   Introduction

The purpose of this project is to create a utility based agent to maximise the number of rows removed in a game of Tetris. This tetris playing agent uses a heuristic function to estimate the utility of each state.

In this report, we discuss the design and strategies used in the agent, and how we got agents to learn better strategies through Genetic Algorithm (Section 2.1, 2.2). We will then look at how we have tackled speeding up gameplays (Section 2.3) and the results of how well our learnt agent performs in a standard game (Section 3). Lastly, we discuss our result and other possible extensions that we can do to make the reduce the time needed for the algorithm to converge.

2   Strategy

The agent's heuristic function sums the linear weights $w(k)$ of features $\varphi_k(s)$ (As stated in Section 2.1) for a given state of the board, s, where n is the number of features as shown below:

$$\hat{V}(s) = \sum_{k=0}^{n} w(k)\varphi_k(s)$$

Where at every turn, the agent evaluates all possible moves and makes the move that gives the best utility.

2.1   Features Selected

This is the list of 11 features that we have selected. They allow us to evaluate each state s based on certain characteristics of the board.

- **NUM_ROWS_REMOVED** – Number of rows removed after each action

- **MAX_HEIGHT** – Height of the tallest column

- **TOTAL_HEIGHT** – Sum of all column heights

- **TOTAL_DIFF_HEIGHT** – Sum of all difference in height of all columns

- **LANDING_HEIGHT** – Height of where the lowest point of next piece lands

- **NUM_HOLES** – Number of empty cells with at least one filled cell above

- **COL_TRANSITION** – Number of filled cells adjacent to empty cells, summed over all columns

- **ROW_TRANSITION** – Same as the above, but applied to rows

- **COVERED_GAPS** – Number of empty cells with a filled cell anywhere above them

- **TOTAL_WELL_DEPTH** – Sum of the depth of all wells

- **HAS_LOST** – Gives a penalty of -10000 if move result in loss, else give 100

Intuitively, the weight vector can be thought of as the agent's strategy on the board, based on the features it sees. For example, if the agent has a weight vector that gives a high positive weight value to TOTAL_DIFF_HEIGHT, it will have a general strategy to always prefer moves that increases the overal "bumpiness" of the tetris wall.

This gives us a clear way in interpreting the values from the weight vector and how it corresponds to the overall strategy of the agent.

2.2 Genetic Algorithm

For our implementation of the genetic algorithm, Each chromosome has a weight vector where each gene (weight value) corresponds to one of the 11 features stated in Section 2.1, and a fitness score.

The fitness score of each chromosome is defined as the mean score of playing 50 games using that individual's chromosome weight.

This is a simple summary of our implementation of the genetic algorithm:

1. Start out with 1000 individuals with random weights, let this be the population pool. Calculate the fitness score of all of the individuals in the population.

2. Select 40% of population via Stochastic Universal Sampling to be potential parents, let this be the parent pool.

3. Generate 40% of population as offsprings and put them into an offspring pool by the process below:

   (a) Randomly select 2 parents from the parent pool generated above

   (b) Crossover with 80% chance, by taking weighted average of genes

   (c) Mutate these 2 offsprings with 8% chance by adding 1/10 times the random gaussian value.

   (d) Calculate fitness score for the 2 offsprings

   (e) Add these 2 offsprings to offspring pool

4. Cull the bottom 40% of the population pool and replace with offsprings in offspring pool

5. Repeat steps 2 to 4 for each generation, till convergence

Convergence is determined by the score of the best individual in the population. If this score has not improved for 50 generations, we terminate the algoritihm.

2.3 Parallelisation and Speedup

Each generation of the algorithm required running games to evaluate fitness. This meant that as the weights progressively got better, each generation started taking a longer time to evaualte.

We decided to parallelise the games by running each game on its own thread. Playing 100 games each, with a set of decent set of weights [1] , the time taken for the parallelised version was 2059 seconds while the sequential version was 6897, giving a speedup of 3.34 times.

Another way that we have tried speeding up the learning algorithm was to reduce the size of the board by reducing its height. Our team ran 2 different learners, one learning on a 9x10 board, while the other learning on a 13x10 board. The learner with 9 rows, even at later generations, took an average of 30 minutes per generation, while the latter, took an average of 1.75 hours per generation.

The machine used for the training and learning the weights, and running all the above tests was SoC compute node xgbp0, it has the following specifications:

2x Intel Xeon Processor E5-2620 v4, 64GB DDR4 RAM, 1x Nvidia Tesla P4 GPU.

3 Results

Figures 1 and 2 shows the fitness scores of the Genetic Algorithm. As can be seen from these graphs, the fitness scores for both tends to plateau at around generation 100, and with the learner in 9 rows converging at generation 464. We can see that the highest score improves less and less as the generations increases. This shows that the algorithm is reaching a maxima.

Figure 3 are from the weights shown in Table 1. These weights were derived from the genetic algorithm learner on a board with 13 rows at generation 132. These were the best set of weights in terms of fitness score from the 13 row learner at the time the 9 row learner reached convergence at generation 464.

The result of running 600 games can be seen in Figure 3, while some common metrics of the 600 games can be seen on Table 2.

---

[1] weight vector used w = [0.00134246, −0.01414993, −0.00659672, 0.00140868, −0.02396361, −0.03055654, −0.06026152, −0.02105507, −0.0340038, −0.0117935, 1] played over 200 games in total, 100 games sequentially and 100 games in parallel, with a total average score of 841279

Figure 1: Fitness scores from learner at 9 rows

Features	Weights
NUM_ROWS_REMOVED	-0.109941154
MAX_HEIGHT	-0.115469783
TOTAL_HEIGHT	-0.043905252
TOTAL_DIFF_HEIGHT	0.0179129081
LANDING_HEIGHT	-0.304447670
NUM_HOLES	-0.386174735
COL_TRANSITION	-0.125186298
ROW_TRANSITION	-0.228061778
COVERED_GAPS	-0.769605890
TOTAL_WELL_DEPTH	-0.193777505
HAS_LOST	0.1367227149

Table 1: Respective Weights for Features

Figure 2: Fitness scores from learner at 13 rows

Figure 3: Results from 600 games

As we can see from the result, most of the games played lie below 30 million lines cleared. However we do see a few outliers that broke 50 million lines cleared, including our best run at 216,319,742 lines cleared.

4   Discussion and Findings

Our group initially decided to use Least Square Policy Iteration (LSPI) by Lagoudakis [4], as our learning algorithm because we found that earlier CS3243 groups such as Tan et al [1] and Nguyen et al [2] had been able to get satisfactory results from it. However, the results that we have produced from our implementation of LSPI was not satisfactory, giving an average of at most 8,000 rows cleared.

We speculated that there may have been an error in our implementation, but despite peer review, we were not able to find the issue. In an attempt to try to improve the learnt weights, we tried implementing other features to better represent the state of the board. This was when our team realised the importance of the choice of features.

We initially implemented our original set of fea-

tures without HAS_LOST. Most of which was taken from previous other works of tetris playing agents, such as the Tetris applications by Colin Fahey and Pierre Dellacherie [3], because of how well those features had worked.

However, those features do not take in account of whether the next move will result in a loss, therefore an important characteristic of the board has not been captured. Implementing HAS_LOST will discourages our agent from making losing moves if there are other available non-losing moves.

After implementing this feature, the number of rows cleared on the average increased from 8,000 to 80,000. Nonetheless, the results were still unsatisfactory and the team decided to change the implementation of the learner to use our own version of genetic algorithm instead.

With regards to the 2 instances of learners running on 9 and 13 rows respectively. We found that the learner on the 9x10 board has converged much faster than the learner on the 13x10 board. However, testing on a full sized 20x10 board with 600

Metrics	Score
Q1 (25th Percentile)	6,307,657.5
Median	13,655,622.0
Q3 (75th Percentile)	25,716,898.5
Mean	19,793,958.2
Max	216,319,742.0
Min	5125.0

Table 2: Common Metrics for the Scores

Metrics	Mean Score
13 Rows Learner (Generation 132)	19,793,958.2
9 Rows Learner (Generation 464)	12,872,842.8

Table 3: Comparison of the final top weights of each learner after 9 rows learner reached convergence

games, the player trained on the larger board gave a higher mean score than the player trained on a smaller board, as can be seen in Table 3.

The apparent tradeoff for this speed up in convergence, is that the set of weights obtained are not ideal for a full sized board, though the player still plays well. We think that learning on a smaller game board may be specialising the player for playing on a smaller game board size but does not enable them to generalise to larger game board sizes.

This illustrates the tradeoff between tweaking the size of the board (to reduce the time taken for convergence), as well as the quality of the weights learnt at the end, which would be another interesting area to look into if we had more time.

5   Further Considerations

In addition to the discussion on Section 2.3, another possible way that we could look into speeding up the fitness evaluation of each chromosome is by reducing the length of each game. If the player has reached the maximum number of moves made or has lost, report the score. However, this may lead to a set of weights that are specialised in playing only up to the maximum number of moves and not anymore. The combination of states that the player sees may also be biased to the states that are closer to the start of a game, thus not giving an evaluation that can be generalised to a standard game. Hence how we tweak the maximum number of moves made would be crucial in balancing the time taken versus quality of weights learnt.

Lastly, we can further parallelise the running of the genetic algorithm by distributing the fitness evaluation of each chromosome to separate machines

in different clusters. This will decrease the computational time needed for each generation as the fitness evaluation for all chromosomes would now be running in parallel in different machines.

6   Conclusion

In this paper, we used Genetic Algorithm to arrive at a strategy for a utility based agent that could play Tetris well. This strategy derived could then be used as good starting point in another algorithm such as LSPI, in order to learn the optimal weights, as our Tetris problem could be tailed to fit a control problem as stated by Lagoudakis [4].

We also looked at different methods of making the algorithm learn a good strategy faster, such as reducing the size of the board, and its various tradeoffs.

As a final concluding point, we think that in order to design a utility based agent that could play Tetris well, we should not rely on only optimising the weights for a set of features, we should also be looking into selecting good features as well. As discussed in Section 4, having a set of features that could better represent the state of the game, results in an agent that plays better. Perhaps another interesting area to research would be looking into creating algorithms that could optimise and learn new features of the board state, and thus the utility function, instead of just optimising its weights.

References

[1] Shawn Tan et al, *Learning about reinforcement learning, with Tetris*, 2014
https://blog.wtf.sg/2014/01/12/learning-about-reinforcement-learning-with-tetris/

[2] Nhan Nguyen et al, *Learning to Play Tetris with Big Data*, 2016
https://github.com/ngthnhan/Tetris/blob/final/repor

[3] Colin P. Fahey, *Tetris AI*, 2003
https://www.colinfahey.com/tetris/tetris.html

[4] Michail G. Lagoudakis, Ronald Parr. *Journal of Machine Learning Research 4 (Dec), 1107-1149*, 2003